

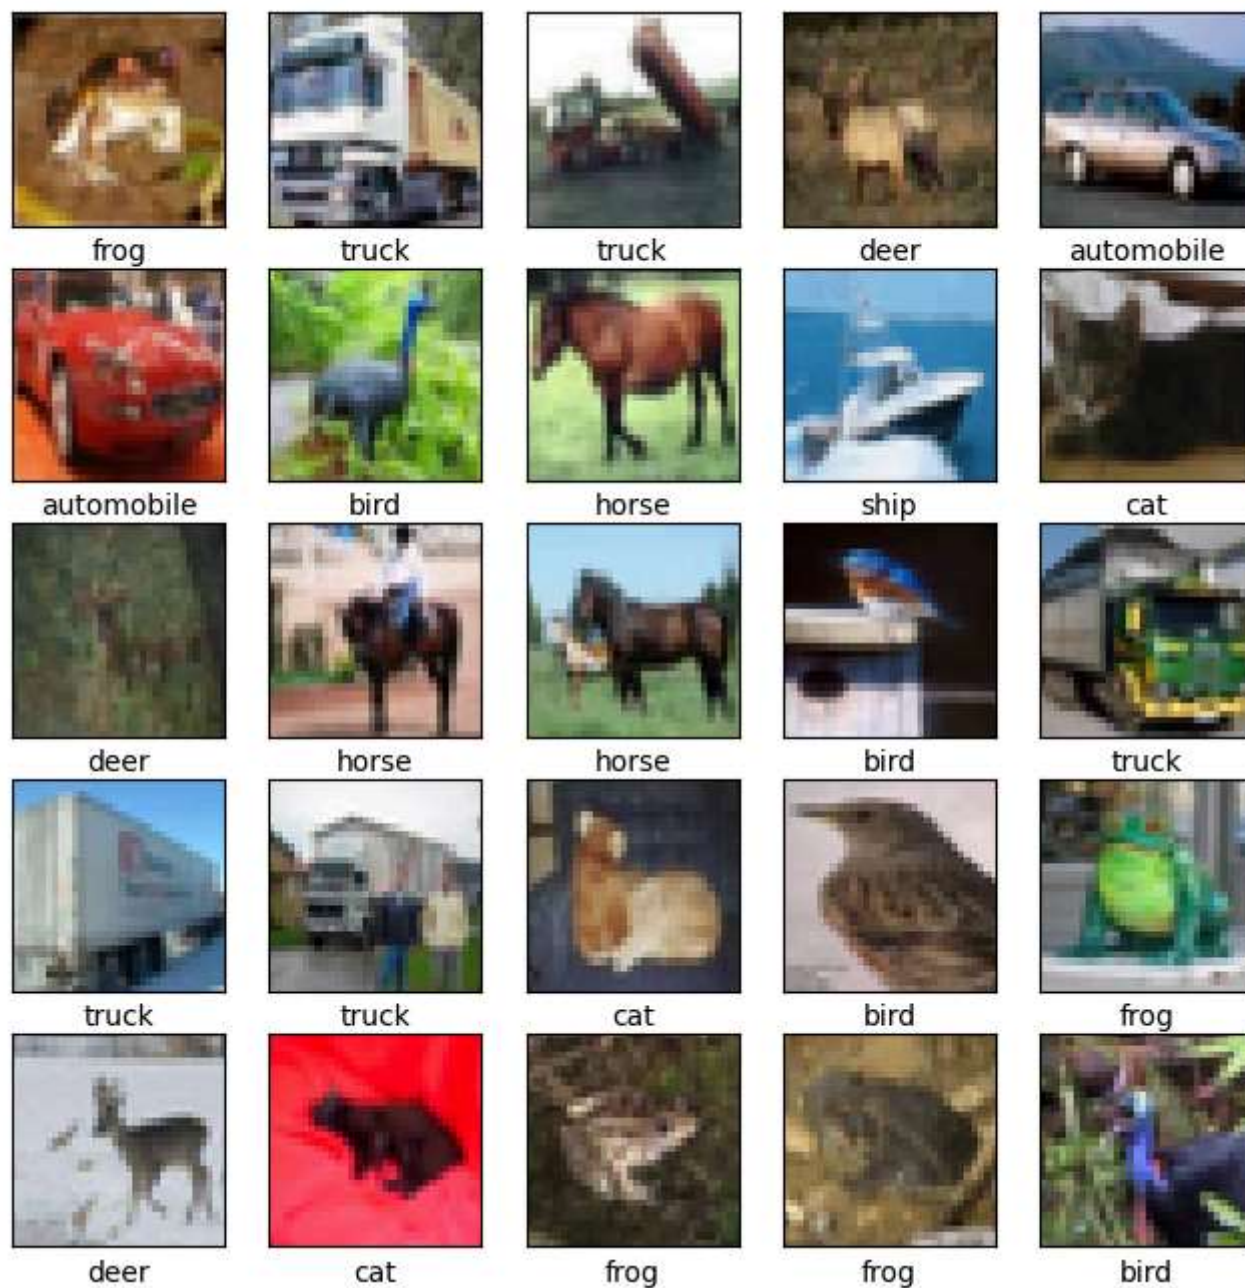
```
In [17]: import tensorflow as tf
        from tensorflow.keras import layers, models
        from tensorflow.keras.datasets import cifar10
        import matplotlib.pyplot as plt

        # Load dataset
        (x_train, y_train), (x_test, y_test) = cifar10.load_data()

        # Normalize pixel values
        x_train, x_test = x_train / 255.0, x_test / 255.0
```

```
In [18]: class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
                        'dog', 'frog', 'horse', 'ship', 'truck']

        plt.figure(figsize=(8,8))
        for i in range(25):
            plt.subplot(5,5,i+1)
            plt.xticks([])
            plt.yticks([])
            plt.grid(False)
            plt.imshow(x_train[i])
            # The CIFAR labels happen to be arrays,
            # which is why we need the extra index
            plt.xlabel(class_names[y_train[i][0]])
        plt.show()
```



```
In [19]: from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
datagen = ImageDataGenerator(  
    rotation_range=15,  
    width_shift_range=0.1,  
    height_shift_range=0.1,  
    horizontal_flip=True,  
)  
datagen.fit(x_train)
```

```
In [20]: model = models.Sequential([  
    layers.Conv2D(32, (3,3), padding='same', activation='relu', input_shape=(32,32,3)),  
    layers.BatchNormalization(),  
    layers.Conv2D(32, (3,3), activation='relu'),  
    layers.BatchNormalization(),  
    layers.MaxPooling2D(pool_size=(2,2)),  
    layers.Dropout(0.25),  
  
    layers.Conv2D(64, (3,3), padding='same', activation='relu'),  
    layers.BatchNormalization(),  
    layers.Conv2D(64, (3,3), activation='relu'),  
    layers.BatchNormalization(),  
    layers.MaxPooling2D(pool_size=(2,2)),  
    layers.Dropout(0.35),  
  
    layers.Flatten(),  
    layers.Dense(128, activation='relu'),  
    layers.BatchNormalization(),  
    layers.Dropout(0.5),  
    layers.Dense(10, activation='softmax')  
)
```

```
In [21]: model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_4 (Conv2D)	(None, 32, 32, 32)	896
batch_normalization_6 (BatchNormalization)	(None, 32, 32, 32)	128
conv2d_5 (Conv2D)	(None, 30, 30, 32)	9,248
batch_normalization_7 (BatchNormalization)	(None, 30, 30, 32)	128
max_pooling2d_2 (MaxPooling2D)	(None, 15, 15, 32)	0
dropout_3 (Dropout)	(None, 15, 15, 32)	0
conv2d_6 (Conv2D)	(None, 15, 15, 64)	18,496
batch_normalization_8 (BatchNormalization)	(None, 15, 15, 64)	256
conv2d_7 (Conv2D)	(None, 13, 13, 64)	36,928
batch_normalization_9 (BatchNormalization)	(None, 13, 13, 64)	256
max_pooling2d_3 (MaxPooling2D)	(None, 6, 6, 64)	0
dropout_4 (Dropout)	(None, 6, 6, 64)	0
flatten_1 (Flatten)	(None, 2304)	0
dense_2 (Dense)	(None, 128)	295,040
batch_normalization_10 (BatchNormalization)	(None, 128)	512
dropout_5 (Dropout)	(None, 128)	0
dense_3 (Dense)	(None, 10)	1,290

Total params: 363,178 (1.39 MB)

Non-trainable params: 640 (2.50 KB)

```
In [23]: from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau

         early_stop = EarlyStopping(patience=10, restore_best_weights=True)
         lr_reduce = ReduceLROnPlateau(monitor='val_accuracy', factor=0.5, patience=3)
```

```
Out[24]: <BatchNormalization name=batch_normalization_11, built=False>
```

[illegible]

```

Epoch 1/10
782/782 ————— 126s 156ms/step - accuracy: 0.3293 - loss: 2.1206 - val_accuracy: 0.5225 - val_loss: 1.3442 - learning_rate: 0.0010
Epoch 2/10
782/782 ————— 113s 144ms/step - accuracy: 0.5225 - loss: 1.3333 - val_accuracy: 0.4927 - val_loss: 1.5585 - learning_rate: 0.0010
Epoch 3/10
782/782 ————— 109s 139ms/step - accuracy: 0.5847 - loss: 1.1737 - val_accuracy: 0.5722 - val_loss: 1.2307 - learning_rate: 0.0010
Epoch 4/10
782/782 ————— 111s 142ms/step - accuracy: 0.6242 - loss: 1.0630 - val_accuracy: 0.6754 - val_loss: 0.9414 - learning_rate: 0.0010
Epoch 5/10
782/782 ————— 110s 141ms/step - accuracy: 0.6529 - loss: 0.9857 - val_accuracy: 0.7186 - val_loss: 0.8221 - learning_rate: 0.0010
Epoch 6/10
782/782 ————— 111s 142ms/step - accuracy: 0.6655 - loss: 0.9510 - val_accuracy: 0.6932 - val_loss: 0.8802 - learning_rate: 0.0010
Epoch 7/10
782/782 ————— 110s 141ms/step - accuracy: 0.6860 - loss: 0.9046 - val_accuracy: 0.6888 - val_loss: 0.9316 - learning_rate: 0.0010
Epoch 8/10
782/782 ————— 112s 143ms/step - accuracy: 0.6924 - loss: 0.8895 - val_accuracy: 0.7153 - val_loss: 0.8368 - learning_rate: 0.0010
Epoch 9/10
782/782 ————— 111s 142ms/step - accuracy: 0.7140 - loss: 0.8232 - val_accuracy: 0.7435 - val_loss: 0.7369 - learning_rate: 5.0000e-04
Epoch 10/10
782/782 ————— 114s 146ms/step - accuracy: 0.7264 - loss: 0.7862 - val_accuracy: 0.7339 - val_loss: 0.7860 - learning_rate: 5.0000e-04

```

```

In [26]: test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)
        print("Test Accuracy:", test_acc)

```

```

313/313 - 4s - 12ms/step - accuracy: 0.7435 - loss: 0.7369
Test Accuracy: 0.7434999942779541

```

```

In [28]: import matplotlib.pyplot as plt

        # Check if 'history' is defined
        if 'history' not in locals():
            raise NameError("The variable 'history' is not defined. Please run the cell where model.fit() is called.")



```

```
# Plot training & validation accuracy and Loss values
plt.figure(figsize=(12, 5))

plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Model Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()

plt.tight_layout()
plt.show()
```



```
In [30]: # Make predictions
predictions = model.predict(x_test)
# Display predictions for the first 5 test images
for i in range(5):
    plt.imshow(x_test[i])
    plt.title(f"Predicted: {class_names[predictions[i].argmax()]}, Actual: {class_names[y_test[i][0]]}")
    plt.axis('off')
    plt.show()
# Save the model
model.save('cifar10_model.h5')
# Load the model (if needed)
# Loaded_model = models.Load_model('cifar10_model.h5')
# Loaded_model.summary()
# Make predictions with the Loaded model
```


Predicted: cat, Actual: cat



Predicted: ship, Actual: ship



Predicted: automobile, Actual: ship



Predicted: automobile, Actual: airplane



Predicted: frog, Actual: frog



WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.